

Convolutional Neural network(CNN) ¶

The convolutional layer is the core building block of a CNN, and it is where the majority of computation occurs. It requires a few components, which are input data, a filter, and a feature map. Let's assume that the input will be a color image, which is made up of a matrix of pixels in 3D. This means that the input will have three dimensions—a height, width, and depth—which correspond to RGB in an image. We also have a feature detector, also known as a kernel or a filter, which will move across the receptive fields of the image, checking if the feature is present. This process is known as a convolution.

```
In [1]: import keras
        from keras.datasets import mnist
        from keras.models import Sequential
        from keras.layers import Dense, Dropout, Flatten
        from keras.layers import Conv2D, MaxPooling2D
        from keras import backend as K
        import numpy as np
```

```
In [2]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>)
11490434/11490434 [=====] - 2s 0us/step

```
In [3]: img_rows, img_cols = 28, 28

        if K.image_data_format() == 'channels_first':
            x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
            x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
            input_shape = (1, img_rows, img_cols)
        else:
            x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
            x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
            input_shape = (img_rows, img_cols, 1)

        x_train = x_train.astype('float32')
        x_test = x_test.astype('float32')
        x_train /= 255
        x_test /= 255

        y_train = keras.utils.to_categorical(y_train, 10)
        y_test = keras.utils.to_categorical(y_test, 10)
```

```
In [5]: model = Sequential()
        model.add(Conv2D(32, kernel_size = (3, 3),
            activation = 'relu', input_shape = input_shape))
        model.add(Conv2D(64, (3, 3), activation = 'relu'))
        model.add(MaxPooling2D(pool_size = (2, 2)))
        model.add(Dropout(0.25)) , model.add(Flatten())
        model.add(Dense(128, activation = 'relu'))
        model.add(Dropout(0.5))
        model.add(Dense(10, activation = 'softmax'))
```

```
In [6]: model.compile(loss = keras.losses.categorical_crossentropy,  
optimizer = keras.optimizers.Adadelta(), metrics = ['accuracy'])
```

```
In [7]: model.fit(  
x_train, y_train,  
batch_size = 128,  
epochs = 12,  
verbose = 1,  
validation_data = (x_test, y_test)  
)
```

Epoch 1/12

469/469 [=====] - 97s 205ms/step - loss: 2.2782 - accuracy: 0.1654 - val_loss: 2.2462 - val_accuracy: 0.3393

Epoch 2/12

469/469 [=====] - 91s 193ms/step - loss: 2.2283 - accuracy: 0.2842 - val_loss: 2.1860 - val_accuracy: 0.5698

Epoch 3/12

469/469 [=====] - 96s 204ms/step - loss: 2.1653 - accuracy: 0.3792 - val_loss: 2.1047 - val_accuracy: 0.6451

Epoch 4/12

469/469 [=====] - 96s 205ms/step - loss: 2.0771 - accuracy: 0.4524 - val_loss: 1.9928 - val_accuracy: 0.6710

Epoch 5/12

469/469 [=====] - 96s 204ms/step - loss: 1.9633 - accuracy: 0.5031 - val_loss: 1.8476 - val_accuracy: 0.7002

Epoch 6/12

469/469 [=====] - 97s 208ms/step - loss: 1.8186 - accuracy: 0.5452 - val_loss: 1.6696 - val_accuracy: 0.7383

Epoch 7/12

469/469 [=====] - 96s 205ms/step - loss: 1.6567 - accuracy: 0.5794 - val_loss: 1.4714 - val_accuracy: 0.7690

Epoch 8/12

469/469 [=====] - 98s 210ms/step - loss: 1.4867 - accuracy: 0.6086 - val_loss: 1.2771 - val_accuracy: 0.7910

Epoch 9/12

469/469 [=====] - 98s 209ms/step - loss: 1.3380 - accuracy: 0.6332 - val_loss: 1.1076 - val_accuracy: 0.8102

Epoch 10/12

469/469 [=====] - 99s 211ms/step - loss: 1.2136 - accuracy: 0.6551 - val_loss: 0.9707 - val_accuracy: 0.8220

Epoch 11/12

469/469 [=====] - 100s 213ms/step - loss: 1.1104 - accuracy: 0.6790 - val_loss: 0.8624 - val_accuracy: 0.8308

Epoch 12/12

469/469 [=====] - 101s 214ms/step - loss: 1.0256 - accuracy: 0.6982 - val_loss: 0.7784 - val_accuracy: 0.8389

Out[7]: <keras.callbacks.History at 0x21d4ce70c40>

```
In [8]: score = model.evaluate(x_test, y_test, verbose = 0)  
  
print('Test loss:', score[0])  
print('Test accuracy:', score[1])
```

Test loss: 0.7784239053726196

Test accuracy: 0.8389000296592712

```
In [9]: pred = model.predict(x_test)
pred = np.argmax(pred, axis = 1)[:5]
label = np.argmax(y_test,axis = 1)[:5]

print(pred)
print(label)
```

```
313/313 [=====] - 4s 12ms/step
[7 2 1 0 4]
[7 2 1 0 4]
```

```
In [ ]:
```